

# Machine vision in Python

Jan Margeta | [jan@kardio.me](mailto:jan@kardio.me) | [@jmargeta](https://twitter.com/@jmargeta)

KardioMe<sup>β</sup>

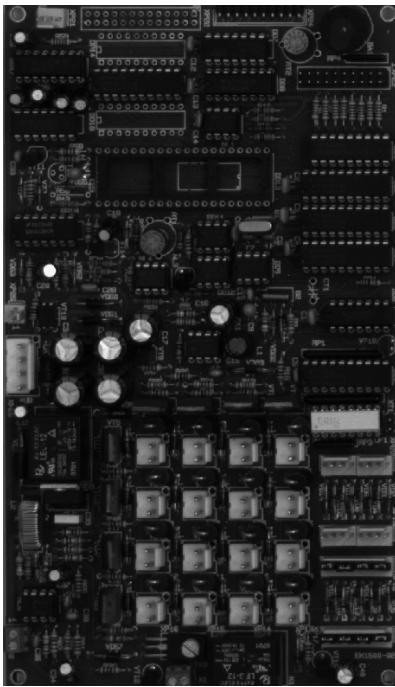
May 4, 2018



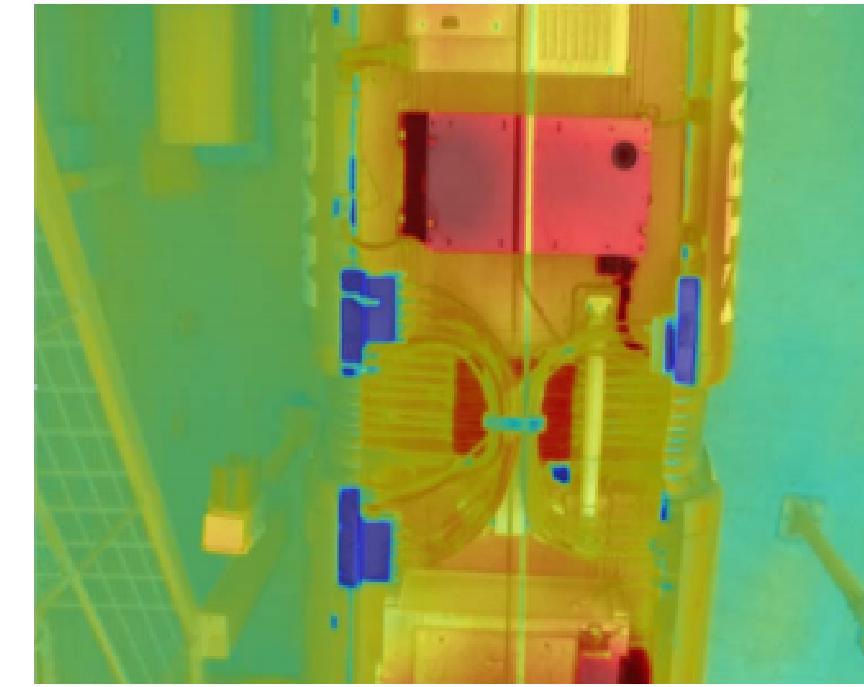
# KardioMe<sup>®</sup>

... ping me if you wanna know more ...

# Machine vision?

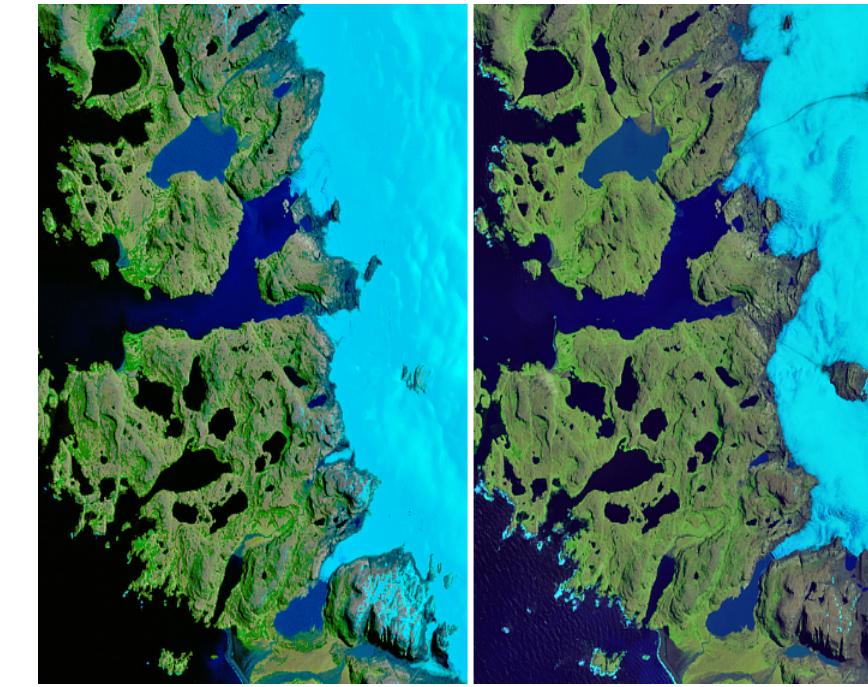


Quality control in  
waste reduction



Early failure  
detection

Courtesy of Alstom transport



Climate change  
monitoring

USGS/NASA Landsat

h more...

A  
n  
d  
m  
u  
c  
h  
m  
u  
c

# Getting some hardware



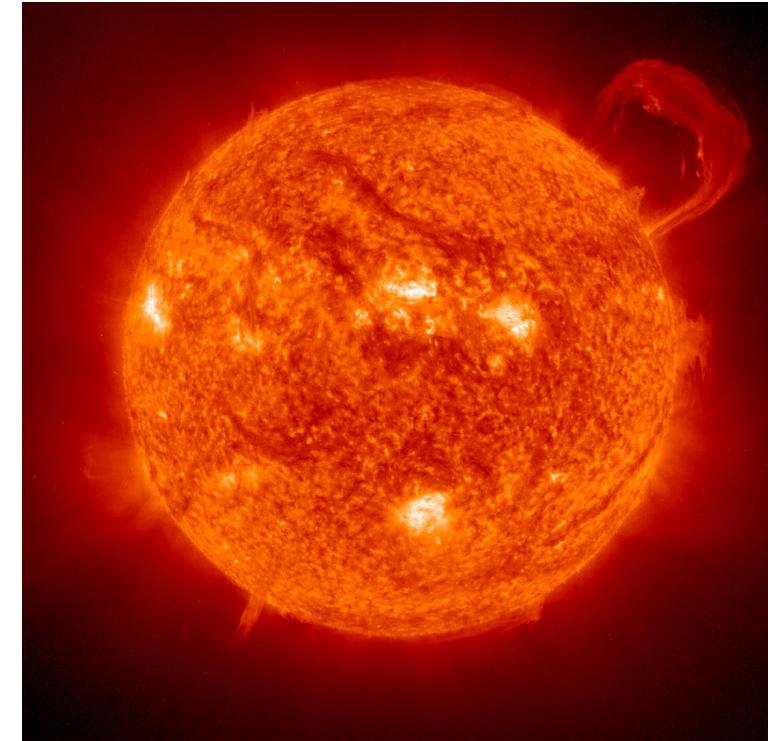
## Camera

resolution  
frame rate  
"color"  
interface



## Optics

field of view  
F-number  
camera compatible



## Lighting

intensity  
strobing  
color  
pattern

# The (usual) Flow

# Camera software

Aravis

Generic

GigE Vision and USB3

C interface

Python via GI

LGPL

Camera vendor's SDK

Vendor lock in

GigE Vision and USB3

Somewhat more reliable

Full cam functionality

C or C++ interfaces only

proprietary

**Let's wrap it!**

# Options for wrapping C and C++ libraries

- Python extension modules
- Swig
- ctypes
- CFFI
- Cython

# Cython I.

## Expose C or C++ structures

```
cdef extern from "CameraParams.h":  
    cppclass GIGE_DEVICE_INFO:  
        unsigned int nIpCfgOption  
        unsigned int nIpCfgCurrent  
        unsigned int nCurrentIp  
        unsigned int nCurrentSubNetMask  
        unsigned int nDefaultGateWay  
        unsigned char chManufacturerName[32]  
        unsigned char chModelName[32]  
        unsigned char chDeviceVersion[32]  
        unsigned char chManufacturerSpecificInfo[48]  
        unsigned char chSerialNumber[16]  
        unsigned char chUserDefinedName[16]  
        unsigned int nNetExport  
        unsigned int nReserved[4]
```

# Cython II.

## Wrap vendor's SDK

```
cdef class GigEDeviceInfo:  
    cdef:  
        GIGE_DEVICE_INFO c_dev_info  
  
    @staticmethod  
    cdef create(GIGE_DEVICE_INFO dev_info):  
        instance = GigEDeviceInfo()  
        instance.c_dev_info = dev_info  
        return instance  
  
    @property  
    def current_ip(self):  
        return self.c_dev_info.nCurrentIp  
    ...
```

# Cython III.

## Compile

```
from distutils.core import setup
from Cython.Build import cythonize

setup(
    name = "Cythonized app",
    ext_modules = cythonize('camera.pyx'),
)
```

# Using the camera



# Find connected cameras

```
import pyhikvision
factory = pyhikvision.camera.TlFactory()
devices = factory.enum_devices()
```

# Setting camera parameters

```
camera = factory.create_device(devices[0])
camera.open(1, 0)
camera.width = 1920
camera.height = 1374
camera.pixel_format = 'Mono8'
camera.frame_rate = 5
camera.exposure_auto = 'continuous'
camera.gain_auto = 'continuous'
```

# Do image acquisition

```
# Allow the camera to start grabbing frames
camera.start_grabbing()

# Grab a frame
frame = camera.grab_frame()

# Async is also (kind of) supported
frame = await camera.grab_frame_async()
```

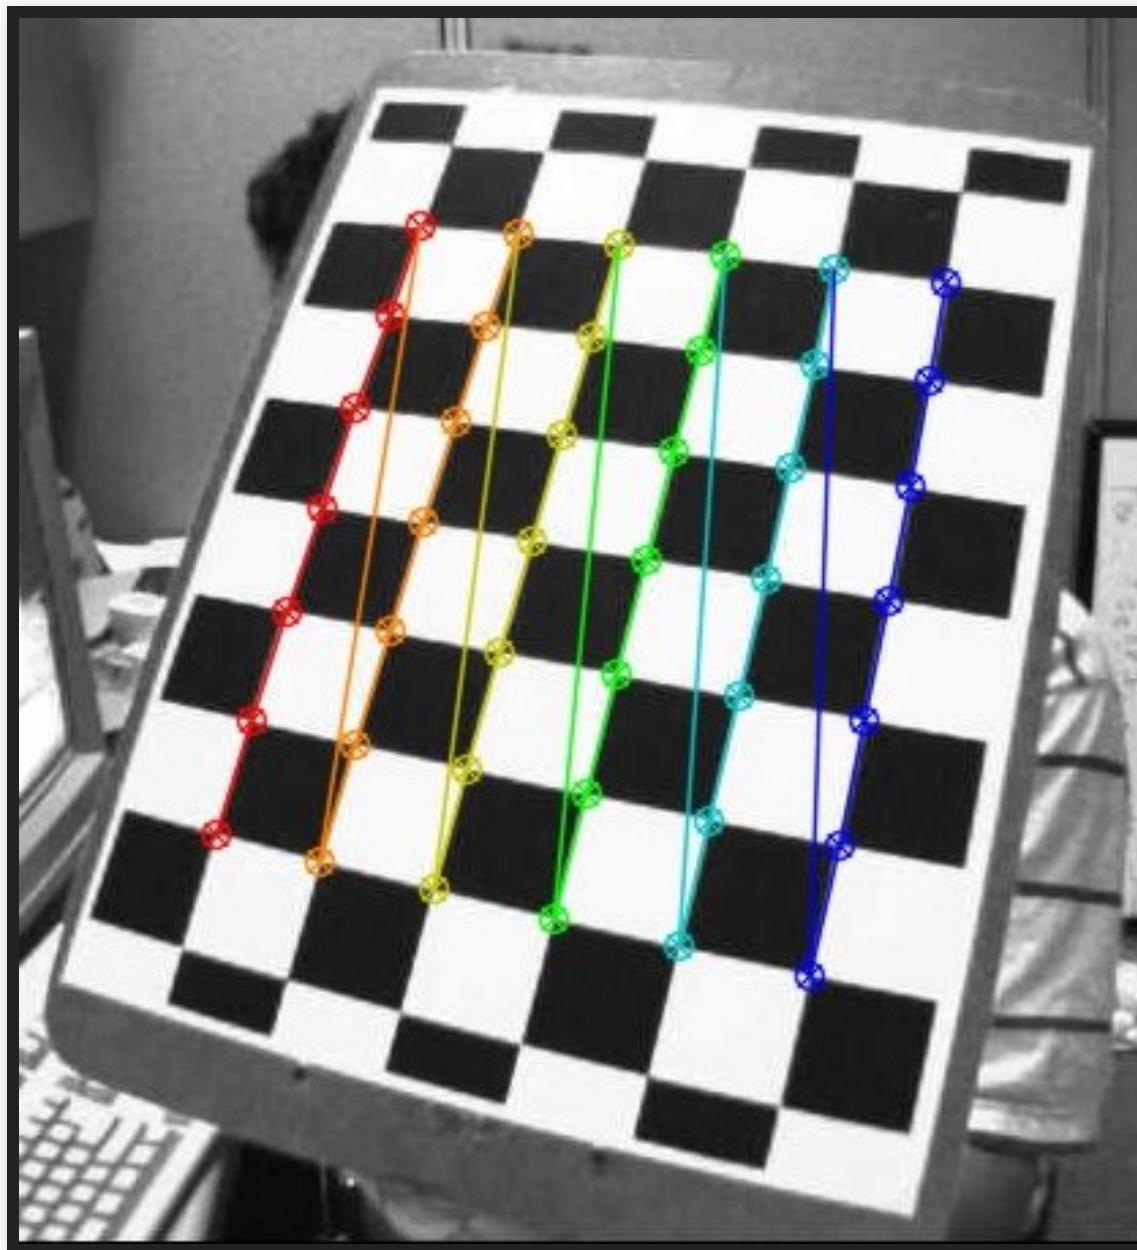
## Frames as numpy arrays!

# Image processing in Python

133	115	57	65	52	69	129	145	134	121	106	81	87	59	42	44	18	24
131	87	53	57	140	143	133	130	119	117	112	115	81	94	48	31	31	25
139	59	54	120	138	135	137	139	120	110	107	109	122	68	105	33	36	16
119	60	46	144	138	138	131	133	120	111	67	45	46	98	91	44	37	18
106	47	71	131	179	155	127	147	111	100	88	88	62	77	43	32	17	
103	43	80	125	127	124	120	109	120	109	106	110	56	83	92	35	34	24
90	51	59	118	123	113	107	111	109	109	99	91	108	103	69	43	31	32
101	60	59	105	105	105	98	99	105	109	95	114	101	110	54	39	36	22
90	45	49	94	101	96	94	109	104	100	107	93	98	39	49	34	31	18
67	46	45	58	95	102	77	85	101	79	91	97	71	58	43	38	16	19
19	39	50	48	54	72	98	91	68	88	97	79	61	36	33	19	17	20

# Camera calibration

Lenses are far from perfect!



Pinhole model

$$K = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

4 projection parameters  
5 distortion parameters

# Camera calibration with OpenCV

```
image_points, object_points = [], []
calibration_object_3d = make_checkerboard_3d(checkerboard_size)

for fname in image_fnames:
    im = load_image(fname)

    board_found, corners = cv2.findChessboardCorners(
        im, checkerboard_size, None)
    if board_found:
        image_points.append(corners_refined)
        object_points.append(calibration_object_3d)

_, camera_matrix, distortion_coeffs, _, _ = cv2.calibrateCamera(
    object_points, image_points, im.shape[::-1], None, None)
```

---

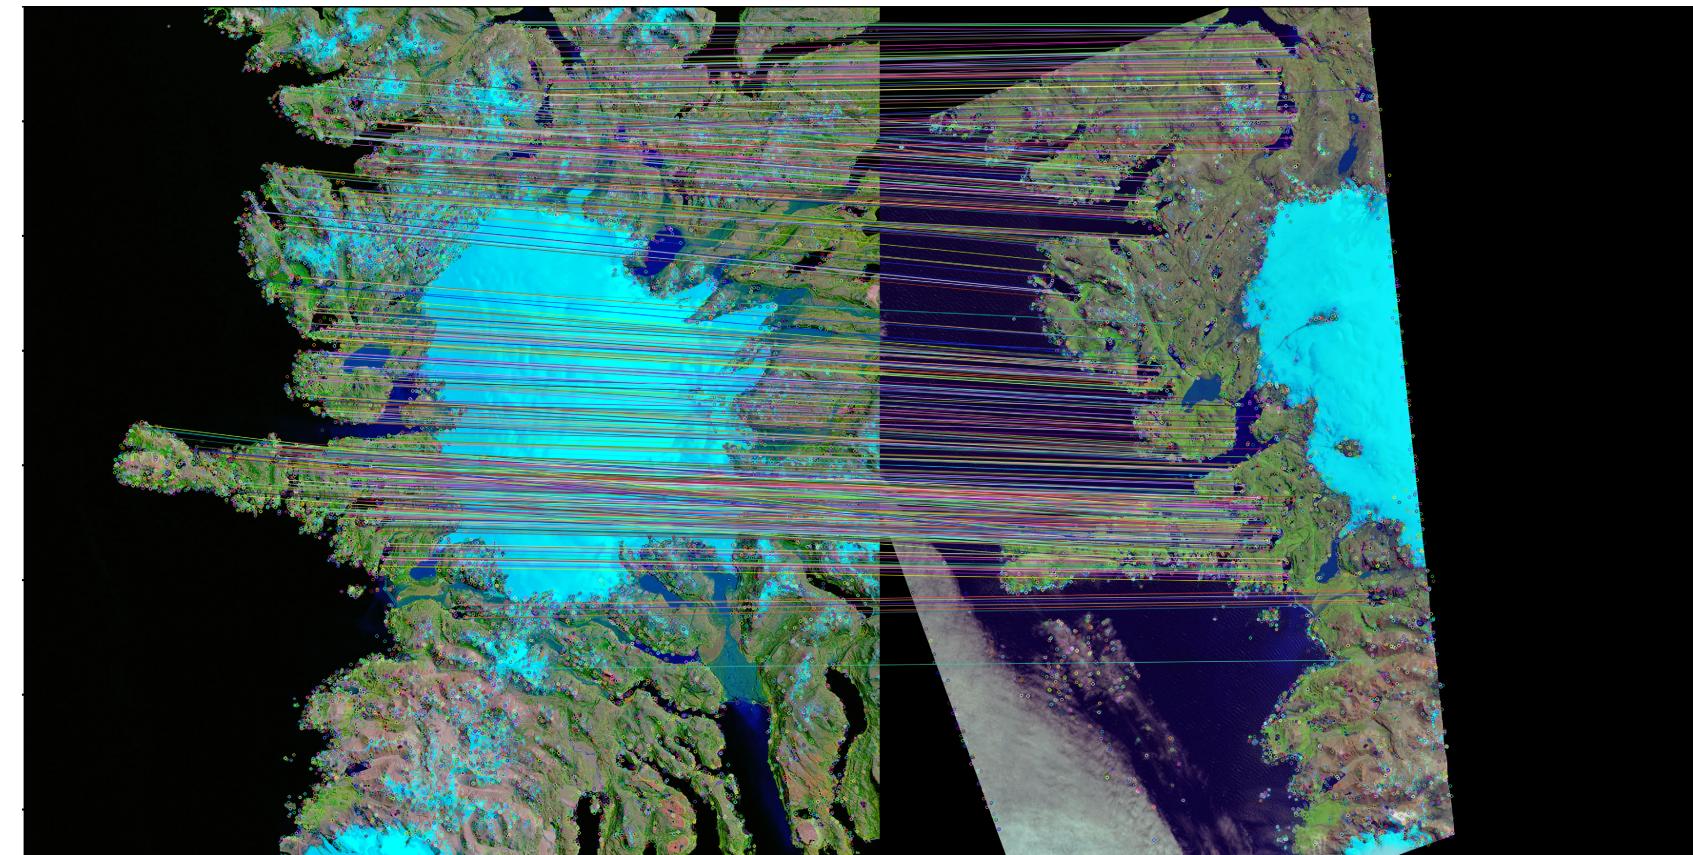
## Undistort images

```
im_undist = cv2.undistort(im, camera_matrix, distortion_coeffs)
```

# Image alignment

600-1000 m glacial retreat over 16 years - Kerguelen Islands - Indian Ocean

November 2001 - January 2017



NASA Earth Observatory images by Joshua Stevens, using Landsat data from the U.S. Geological Survey

# Finding image correspondances

```
import cv2

# find the keypoints and descriptors with ORB
features = cv2.AKAZE_create()
keypoints_ref, descriptors_ref = features.detectAndCompute(im_ref, None)
keypoints, descriptors = features.detectAndCompute(im, None)

# create BFMatcher object
bf = cv2.BFMatcher(cv2.NORM_HAMMING, crossCheck=True)

# Match descriptors.
matches = bf.match(descriptors_ref, descriptors)

# Sort them in the order of their distance.
matches = sorted(matches, key=lambda x:x.distance)
```

# Robust image alignment - RANSAC

```
from skimage import transform
from skimage.measure import ransac

# transformation of matched keypoints to their coordinates
# (points, points_ref) omitted for brevity

model_robust, inliers = ransac(
    (points, points_ref),
    transform.SimilarityTransform, min_samples=3,
    residual_threshold=2, max_trials=100)
```

# Image alignment

**600-1000 m glacial retreat over 16 years - Kerguelen Islands - Indian Ocean**

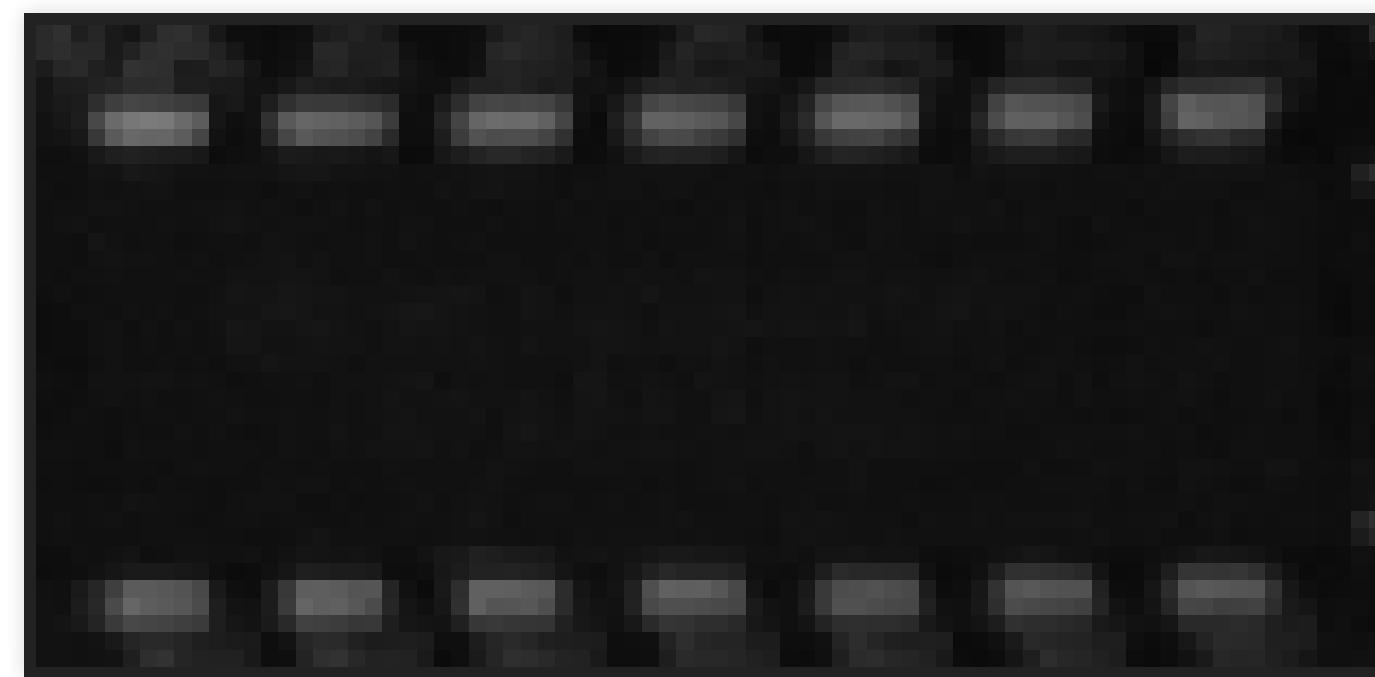
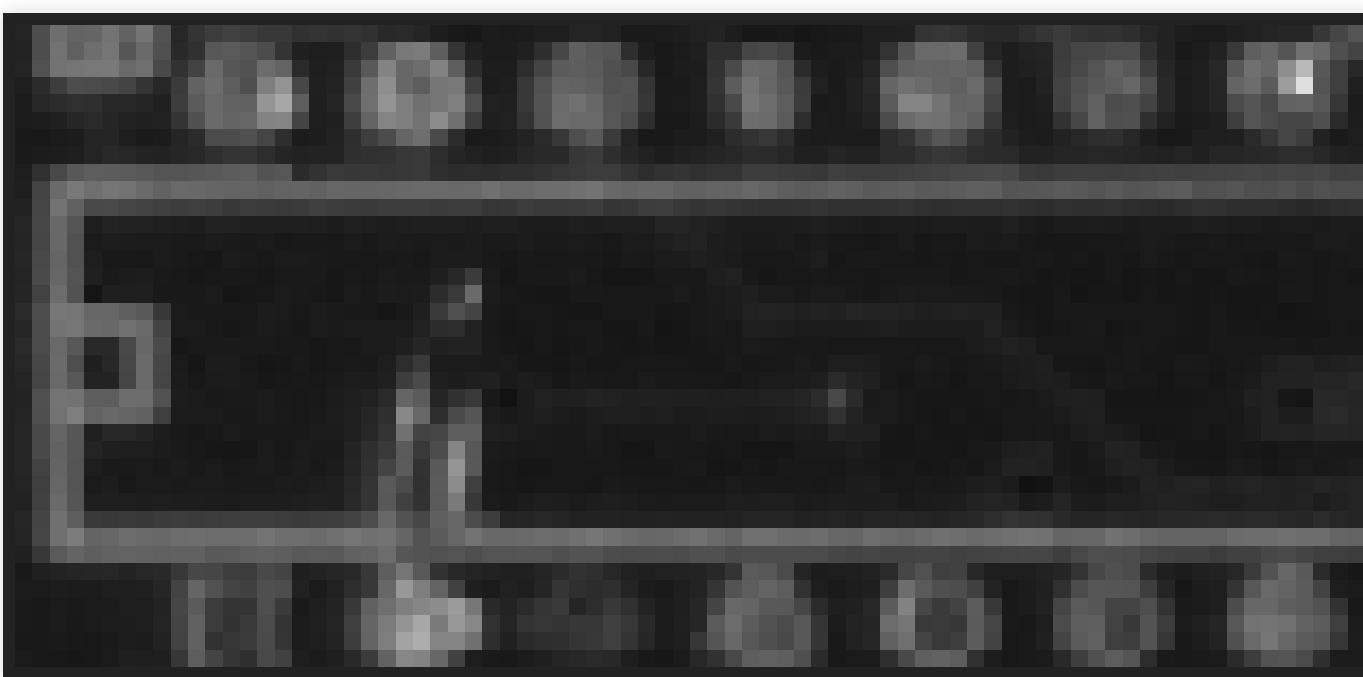
NASA Earth Observatory images by Joshua Stevens, using Landsat data from the U.S. Geological Survey



Mosaics - for tram  
maintenance

# Object detection

Image -> Yes/No



Writing small reusable  
modules for computer  
vision in Gluon

# Example detection model in Gluon

```
def make_detection_model(num_output_classes=4):  
    net = gluon.nn.Sequential()  
    net.add(nn.Conv2d(10), activation='relu')  
    net.add(nn.GlobalAvgPool2D())  
    net.add(nn.Flatten())  
    net.add(nn.Dense(2))  
    return net
```

# Training a module

```
for batch in train_data:  
    data = batch.data[0].as_in_context(computation_context)  
    label = batch.label[0].as_in_context(computation_context)  
  
    # record the computational graph  
    with mx.autograd.record():  
        prediction = net(data)  
  
        # compute the loss  
        loss = loss_function(prediction, label)  
  
        # compute the gradients  
        loss.backward()  
  
        # run one optimization step  
        trainer.step(data.shape[0])
```

Easy to switch the training on and off

# Async

## Acquisition, processing, and real-time updates with Sanic

```
from sanic import Sanic
app = Sanic(__name__)
...
app.add_task(main())
app.run(host="0.0.0.0", port=8000, debug=True)
```

```
async def main():
    camera_manager = app.camera_manager

    acquisition_queue = Queue(maxsize=1)

    app.loop.create_task(image_producer(
        camera_manager=camera_manager,
        queue=acquisition_queue
    ))
    ...

```

# Producers and consumers

```
async def image_producer(camera_manager, queue):
    logger.debug(f'Starting image producer')
    while True:
        cameras = camera_manager.cameras
        logger.debug(f'Found {len(cameras)} cameras')

        await asyncio.gather(
            *[grab_and_enqueue(cam, queue) for cam in cameras],
            return_exceptions=True
        )
        await queue.put(None)
```

# Real-time GUI updates

```
class StateHub:  
    def __init__(self, serializer: Callable):  
        self.serializer = serializer  
        self.subscribers: set = set()  
        self.cache: Any = None  
  
    @async def subscribe(self, subscriber, replay_last=True):  
        self.subscribers.add(subscriber)  
        if replay_last and self.cache is not None:  
            await subscriber.send(self.cache)  
  
    def unsubscribe(self, subscriber):  
        self.subscribers.remove(subscriber)
```

# Publish updates

```
async def send(self, message) -> None:  
    """Publish the message to all subscribers."""  
    if self.serializer:  
        message = self.serializer(message)  
  
    self.cache = message  
  
    subscribers = self.subscribers.copy()  
    for subscriber in subscribers:  
        try:  
            await subscriber.send(message)  
        except ConnectionClosed:  
            self.unsubscribe(subscriber)
```

# Prepare the state update publisher

```
@app.listener('before_server_start')
async def prepare_info_hub(app, loop):
    """Prepare the hub"""
    app.hub = StateHub(serializer=serializer)
```

# Websockets for real-time updates

```
@app.websocket('/feed')
async def feed_camera(request, ws):
    await app.hub.subscribe(ws)
    # prevent the websocket from closing
    while True:
        await asyncio.sleep(0.001)
```

# Deployment

- Camera drivers can be nasty
- Docker and friends make things simpler to deploy and restart

# Conclusion

- Vision in python is fun and a pleasure to work with
- OpenCV, scikit-image, imageio,...
- Python is fast enough
- 2-3 frames per second (12MPix images)

# Machine vision in Python

Jan Margeta | [jan@kardio.me](mailto:jan@kardio.me) | [@jmargeta](https://twitter.com/@jmargeta)

KardioMe<sup>β</sup>

May 4, 2018